

Asymmetric Trust in Distributed Systems

Luca Zanolini
University of Bern

Supervisor: Prof. Dr. **Christian Cachin**

6th July 2023

Secure distributed systems rely on **trust**

- ◆ Specifies the **failures** that a system can tolerate.
- ◆ Determines the **conditions** under which a system operates correctly.
- ◆ Defined through a **fail-prone system**.
- ◆ **Fail-prone systems** are useful tools for the design of distributed algorithms.



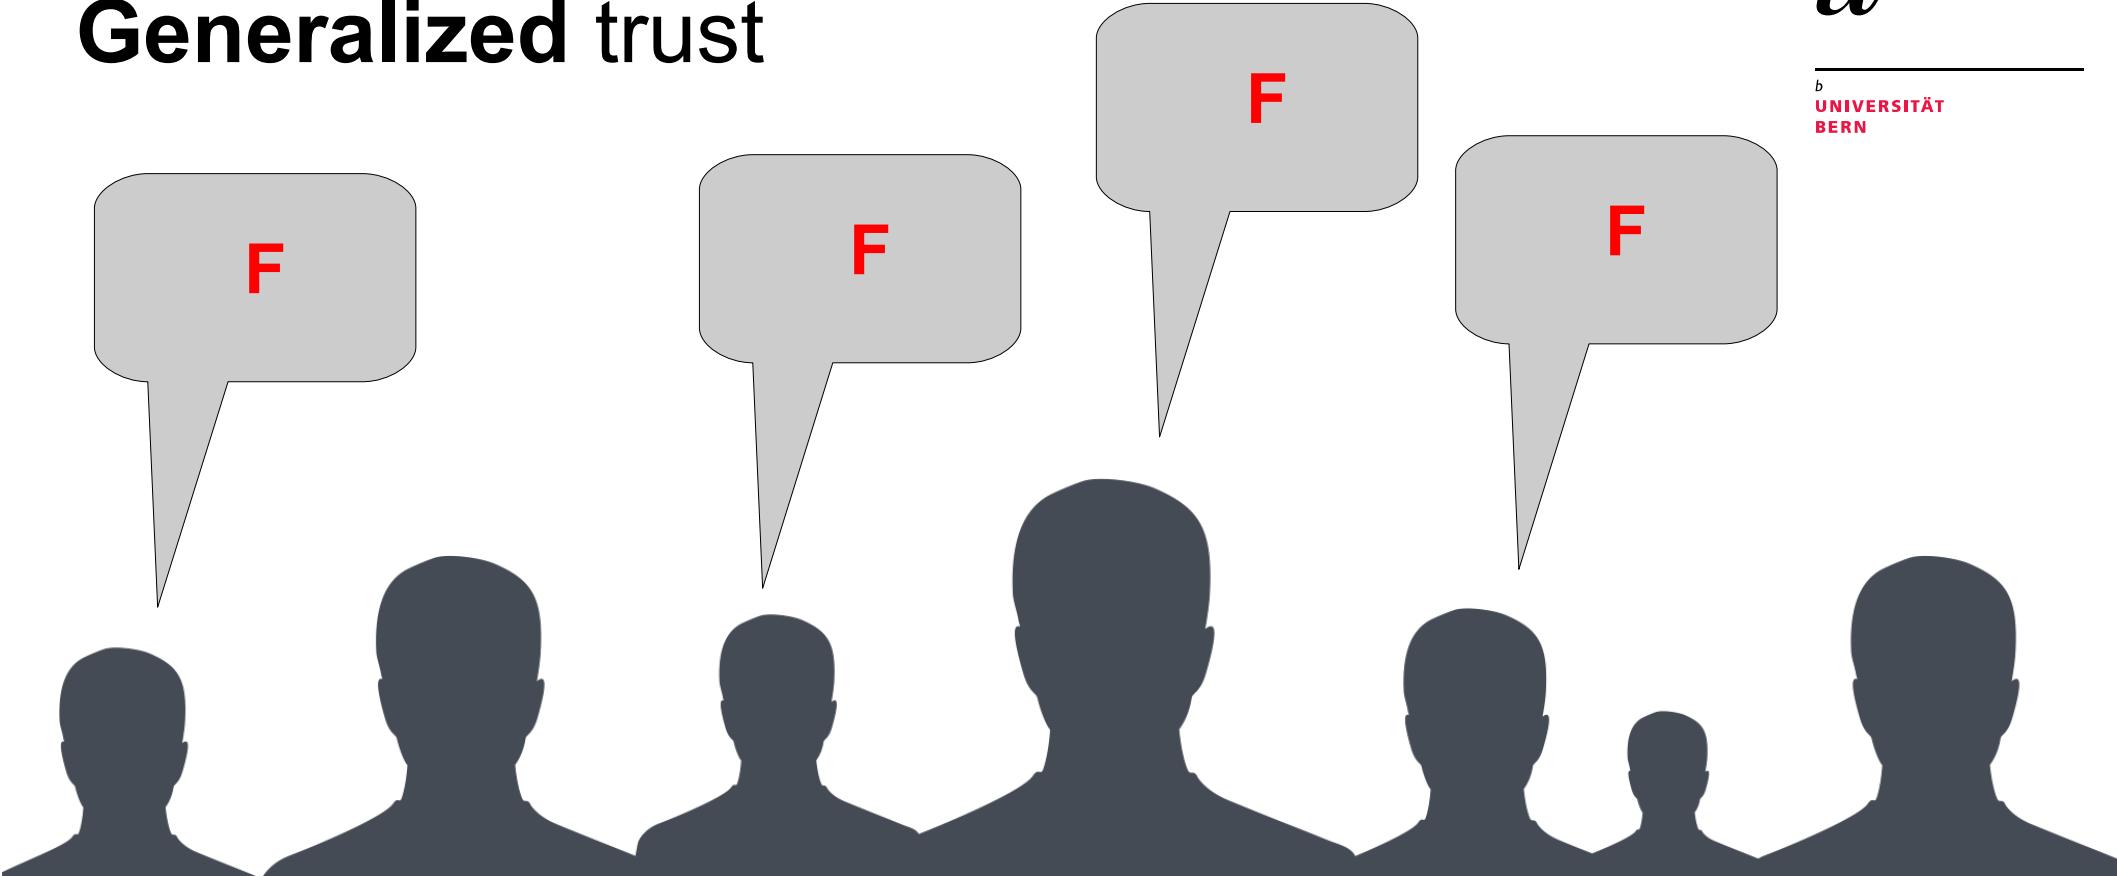
Permissioned systems

- ◆ $P = \{p_1, \dots, p_n\}$.
- ◆ Full system membership is public knowledge.
- ◆ Trust assumptions are public knowledge.
- ◆ Participants do not lie about their trust assumptions.

u^b

^b
UNIVERSITÄT
BERN

Generalized trust

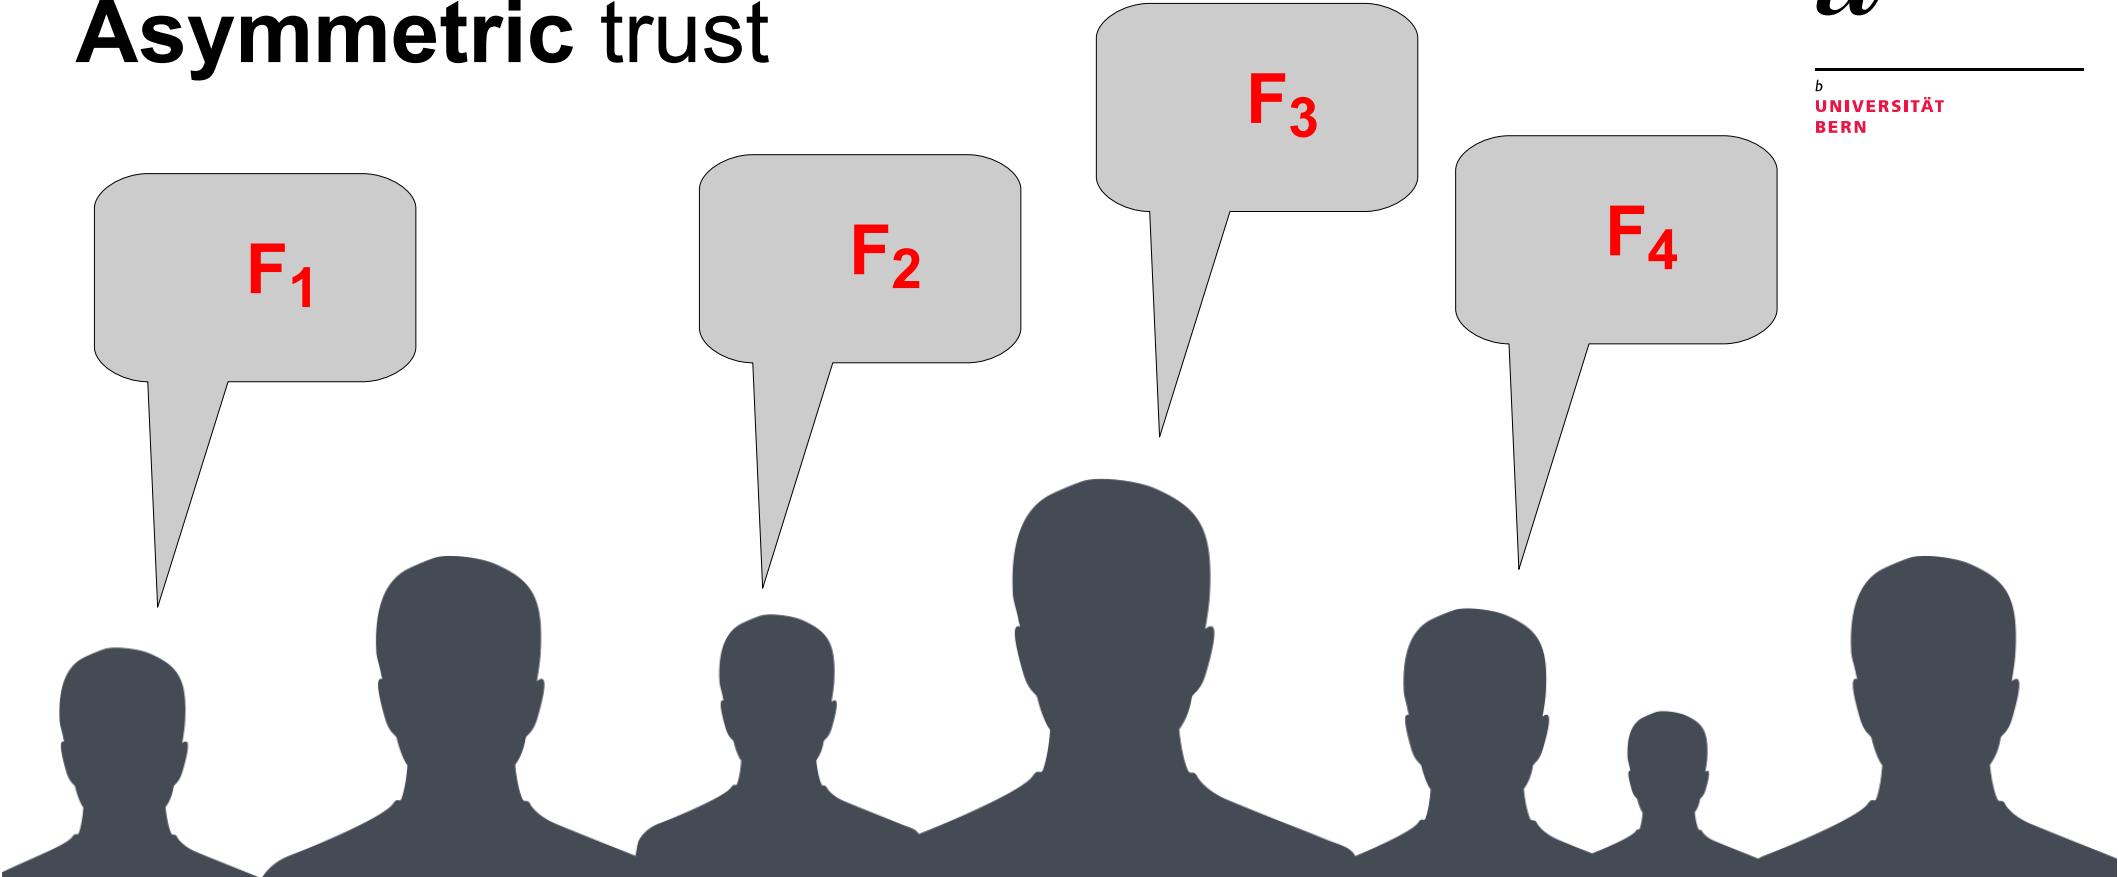


Byzantine quorum systems

- ◆ Set of processes $\mathbf{P} = \{p_1, \dots, p_n\}$.
- ◆ Fail-prone system $\mathbf{F} \subseteq 2^{\mathbf{P}}$: all processes in some $F \in \mathbf{F}$ may fail together.
- ◆ Quorum system $\mathbf{Q} \subseteq 2^{\mathbf{P}}$, where any $Q \in \mathbf{Q}$ is a **quorum**, if and only if:
 - **Consistency**:
$$\forall Q_1, Q_2 \in \mathbf{Q}, \forall F \in \mathbf{F} : Q_1 \cap Q_2 \not\subseteq F.$$
 - **Availability**:
$$\forall F \in \mathbf{F} : \exists Q \in \mathbf{Q} : F \cap Q = \emptyset.$$

[Malkhi & Reiter, 1998]

Asymmetric trust



Asymmetric Byzantine quorum systems

- ◆ Set of processes $\mathbf{P} = \{p_1, \dots, p_n\}$.
- ◆ Fail-prone systems $F_i \subseteq 2^P$ for p_i .
- ◆ Quorum systems $Q_i \subseteq 2^P$, where any $Q_i \in Q_i$ is a **quorum** for p_i , if and only if:
 - **Consistency:**
 $\forall Q_i \in Q_i, \forall Q_j \in Q_j, \forall F_{ij} \in F_i^* \cap F_j^*: Q_i \cap Q_j \not\subseteq F_{ij}$.
 - **Availability:**
 $\forall F_i \in F_i: \exists Q_i \in Q_i: F_i \cap Q_i = \emptyset$.

[Cachin & Tackmann, 2019]

In the **asymmetric** trust model

- ◆ **Faulty:** A process $p_i \in F$ is called *faulty*
- ◆ **Naive:** A correct process p_i for which $F \notin F_i^*$ is called *naive*
- ◆ **Wise:** A correct process for which $F \in F_i^*$ is called *wise*

Guild



u^b

^b
UNIVERSITÄT
BERN

Some of our results

- ◆ Better understanding of the **relationship** between wise and naive processes.
- ◆ **Uniqueness** of the guild in an execution.
- ◆ Importance of a guild in **kernel-based protocols**, e.g., Bracha broadcast.
- ◆ Tolerated system $T = \{P \setminus G, \text{ for any possible guild } G \}$
- ◆ **Composition rule** among asymmetric-trust based systems.

Find a (deterministic) composition rule

F_1

F_2

F_3

F_4

F_5

F_6

F_7

F_8

F_5

F_9

[$F_1, F_2, F_3, F_4, F_5, F_6, F_7, F_8, F_9$]

First asynchronous Byzantine **consensus** **protocol** with asymmetric trust

u^b

^b
UNIVERSITÄT
BERN

- ◆ It uses **randomization**
- ◆ Signature-free
- ◆ Round-based
- ◆ Suitable for applications in **blockchain** networks
- ◆ Builds on the protocol by Mostéfaoui et al. (PODC 2014)

Signature-Free Asynchronous Byzantine Consensus with $t < n/3$ and $O(n^2)$ Messages

Achour Mostéfaoui
LINA, Université de Nantes
44322 Nantes, France
Achour.Mostefaoui@univ-nantes.fr

Hamouma Moumen
University of Bejaia
Bejaia, Algeria
moumenh@gmail.com

Michel Raynal
Institut Universitaire de France
& IRISA, Université de Rennes
raynal@irisa.fr

The (original) protocol

- i. Binary validated **broadcast**
- ii. Randomized **consensus**
 - Uses a *common coin*

bv-broadcast(b)

The (original) protocol

- i. Binary validated **broadcast**
- ii. Randomized **consensus**
 - Uses a *common coin*

$bv\text{-broadcast}(b) \rightarrow bv\text{-deliver}(b)$
 $2f+1$

The (original) protocol

- i. Binary validated **broadcast**
- ii. Randomized **consensus**
 - Uses a *common coin*

$bv\text{-broadcast}(b) \rightarrow bv\text{-deliver}(b) \rightarrow [AUX, b]$ to all
 $2f+1$

The (original) protocol

- i. Binary validated **broadcast**
- ii. Randomized **consensus**
 - Uses a *common coin*

$bv\text{-broadcast}(b) \rightarrow bv\text{-deliver}(b) \rightarrow [AUX, b] \text{ to all} \rightarrow b \text{ received}$
 $2f+1$ $2f+1$

The (original) protocol

- i. Binary validated **broadcast**
- ii. Randomized **consensus**
 - Uses a ***common coin***

$bv\text{-broadcast}(b) \rightarrow bv\text{-deliver}(b)$ → [AUX, b] to all → b received → release-coin
 $2f+1$ $2f+1$

The (original) protocol

- i. Binary validated **broadcast**
- ii. Randomized **consensus**
 - Uses a ***common coin***

$bv\text{-broadcast}(b) \rightarrow bv\text{-deliver}(b)$ → [AUX, b] to all → b received → release-coin → output-coin(s)
 $2f+1$ $2f+1$

The (original) protocol

- i. Binary validated **broadcast**
- ii. Randomized **consensus**
 - Uses a *common coin*

$bv\text{-broadcast}(b) \rightarrow bv\text{-deliver}(b) \rightarrow [AUX, b] \text{ to all} \rightarrow b \text{ received} \rightarrow release\text{-coin} \rightarrow output\text{-coin}(s) \rightarrow$

$if b = s, rbc\text{-decide}(b)$	$if b \neq s, bv\text{-broadcast}(b)$	$if \{0, 1\}, bv\text{-broadcast}(s)$
$2f+1$	$2f+1$	

\uparrow

B

Liveness issue!

The network **reorders** messages between correct processes and delays them until the coin value becomes known.

Fixing the problem

- i. **FIFO** ordering on the reliable point-to-point links, including the messages exchanged by the coin implementation
 - the adversary may no longer exploit its knowledge of the coin value to prevent termination.
- ii. Allow the set B to **dynamically change** while the coin protocol executes.
- iii. Our protocol **does not** execute rounds forever, as in the original formulation.

The (asymmetric) protocol

- i. Asymmetric binary validated **broadcast**
- ii. Asymmetric randomized **consensus**
 - Uses an **asymmetric common coin**

abv-broadcast(b) → abv-deliver(b) → [AUX,b] to all → b received → release-coin → output-coin(s) → if b = s, arbc-decide(b)

Q_i

Q_j

Asymmetric strong Byzantine consensus

In all executions with a guild:

- ◆ **[Probabilistic termination]** Every wise process decides with probability 1.
- ◆ **[Strong validity]** A wise process only decides a value that has been proposed by some processes in the maximal guild.
- ◆ **[Integrity]** No correct process decides twice.
- ◆ **[Agreement]** No two wise processes decide differently.

Permissionless systems

- ◆ $\mathbf{P} = \{p_1, p_2, \dots\}$.
- ◆ Knowledge of the full system membership is **not** available.
- ◆ Trust assumptions are (**partially**) public knowledge.
- ◆ Participants **can lie** about their trust assumptions.

Our model

- ◆ Each process p_i makes assumptions about a set $P_i \subseteq P = \{p_1, p_2, \dots\}$ called p_i 's trusted set, using a fail-prone system F_i over P_i .
- ◆ Point-to-point communication & best-effort gossip primitive.
- ◆ Each process p_i continuously discovers new processes and learns their assumptions.
- ◆ A **permissionless fail-prone system** is an array:

$$F = [(P_1, F_1), (P_2, F_2), \dots,]$$

Assumptions of p_i

We say that the assumptions of a process p_i are satisfied in an execution if the set A of processes that fail is such that there exists a fail-prone set $F \in F_i$ such that:

- i. $A \cap P_i \subseteq F$;
- ii. the assumptions of every member of $P_i \setminus F$ are satisfied.

Assumptions of p_i

We say that the assumptions of a process p_i are satisfied in an execution if the set A of processes that fail is such that there exists a fail-prone set $F \in F_i$ such that:

- i. $A \cap P_i \subseteq F$;
- ii. the assumptions of every member of $P_i \setminus F$ are satisfied.

If p_i has its assumptions satisfied in an execution, we say that p_i tolerates the execution.

A set of processes L tolerates a set of processes A if and only if every process p_i in $L \setminus A$ tolerates an execution with set of faulty processes A .

A new kind of failure assumptions

A participant's assumption are not only about failures, but *also about whether other participants make correct assumptions.*

A new kind of failure assumptions

A participant's assumption are not only about failures, but *also about whether other participants make correct assumptions.*

How do we define quorums?

- ◆ Global **intersection** property among quorums?
- ◆ Malicious processes can **lie** about their assumptions.

Views

A **view** $V = [V_1, V_2, \dots]$ is an array with one entry $V[j] = V_j$ for each process p_i such that:

- i. either V_j is the special value \perp ; or
- ii. $V_j = (P_j, F_j)$ consists of a set of processes P_j and a fail-prone system F_j .

A process p_i 's view is what p_i thinks other's assumptions are. However, such view might contain lies from Byzantine processes.

Views

A **view** $\mathbb{V} = [\mathbb{V}_1, \mathbb{V}_2, \dots]$ is an array with one entry $\mathbb{V}[j] = \mathbb{V}_j$ for each process p_i such that:

- i. either \mathbb{V}_j is the special value \perp ; or
- ii. $\mathbb{V}_j = (P_j, F_j)$ consists of a set of processes P_j and a fail-prone system F_j .

A process p_i 's view is what p_i thinks other's assumptions are. However, such view might contain lies from Byzantine processes.

Given a set of faulty processes A in an execution, we say that a view \mathbb{V} is **A-resilient** if and only if for every process $p_i \notin A$, either $\mathbb{V}[i] = \perp$ or $\mathbb{V}[i] = F[i]$.

Quorum function

A quorum is a set of processes that satisfies the assumptions of every one of its members.

Quorum function

A quorum is a set of processes that satisfies the assumptions of every one of its members.

The **quorum function** $Q : P \times V \rightarrow 2^P$ maps a process p_i and a view V to a set of processes such that $Q \in Q(p_i, V)$ if and only if:

- i. there exists $F \in F_i$ for p_i such that $P_i \setminus F \subseteq Q$;
- ii. for every process $p_j \neq p_i \in Q$ with $V[i] \neq \perp$ and $V_j = (P_j, F_j)$, there exists $F \in F_j$ for p_j such that $P_j \setminus F \subseteq Q$.

Every element of $Q \in Q(p_i, V)$ is called a **permissionless quorum** for p_i .

Leagues

A league is a set of processes that enjoys quorum intersection and quorum availability in all executions that it tolerates.

Leagues

A *league* is a set of processes that enjoys quorum intersection and quorum availability in all executions that it tolerates.

A set of processes L is a **league** for the quorum function Q , if and only if:

- i. **Consistency:** for every set $A \subseteq P$ tolerated by L , for every two A -resilient views V and V' , for every two processes p_i and $p_j \in L \setminus A$, and for every two quorums $Q_i \in Q(p_i, V)$ and $Q_j \in Q(p_j, V')$ it holds $(Q_i \cap Q_j) \setminus A \neq \emptyset$;
- ii. **Availability:** for every set $A \subseteq P$ tolerated by L and for every process $p_i \in L \setminus A$, there exists a quorum $Q_i \in Q(p_i, F)$ such that $Q_i \subseteq L \setminus A$.

Permissionless Byz. reliable broadcast

For every league L and every execution tolerated by L :

- **[Validity]** If a correct process p_s broadcasts a value v , the all correct processes in L eventually deliver v .
- **[Integrity]** For any value v , every correct process delivers v at most once. Moreover, if the sender p_s is correct and the receiver is correct and in L , then v was previously broadcast by p_s .
- **[Consistency]** If a correct process in L delivers some value v and another correct process in L delivers some value v' , then $v = v'$.
- **[Totality]** If a correct process in L delivers some value v , then all correct processes in L eventually deliver some value.

Open questions

- ◆ Asymmetric threshold cryptography.
- ◆ Asymmetric leader-based consensus protocols.
- ◆ More composition rules.
- ◆ Byzantine consensus protocols in the permissionless setting.

Thank you!

Bibliography

- ◆ O. Alpos, C. Cachin, and **L. Zanolini**, “How to trust strangers: Composition of byzantine quorum systems,” in **SRDS** 2021
- ◆ C. Cachin, G. Losa, and **L. Zanolini**, “Quorum systems in per- missionless networks,” in **OPODIS** 2022
- ◆ C. Cachin and **L. Zanolini**, “Asymmetric asynchronous byzantine consensus,” in **CBT 2021 - ESORICS** 2021
- ◆ C. Cachin and **L. Zanolini**, “Brief announcement: Revisiting signature-free asynchronous byzantine consensus,” in **DISC** 2021

Asymmetric common coin

A protocol for asymmetric common coin satisfies the following properties:

- **[Termination]** In all the executions with a guild, every process in the maximal guild eventually outputs a coin value.
- **[Unpredictability]** In all the executions with a guild, no process has any information about the value of the coin before at least a kernel for all wise processes, which consists of correct processes, has released the coin.
- **[Matching]** In all the executions with a guild, with probability 1 every process in the maximal guild outputs the same coin value.
- **[No bias]** The distribution of the coin is uniform over $\{0,1\}$.

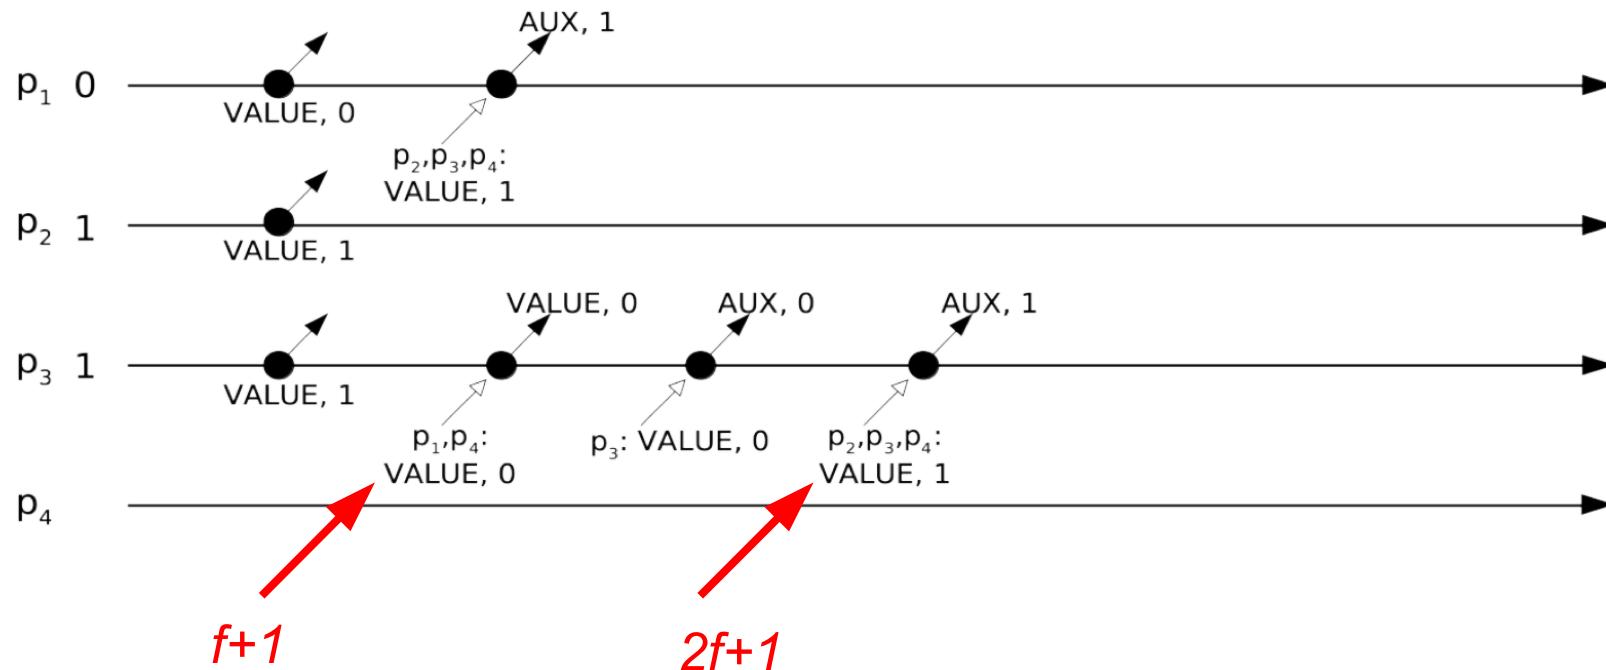
Algorithm 2 Asymmetric common coin for round *round* (code for p_i)

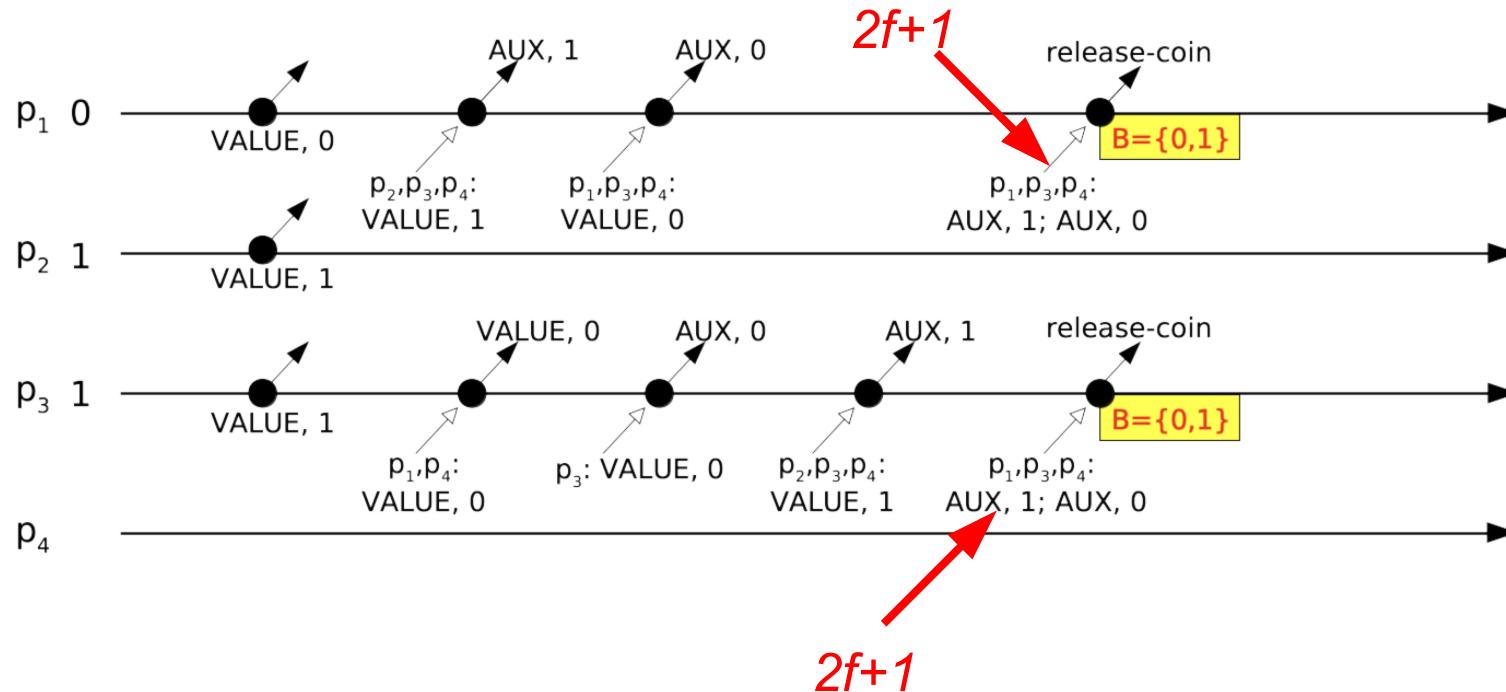
```
1: State
2:    $\mathcal{H}$ : set of all possible guilds
3:    $share[\mathcal{G}][j]$ : if  $p_i \in \mathcal{G}$ , this holds the share received from  $p_j$ 
4:     for guild  $\mathcal{G}$ ; initially  $\perp$ 

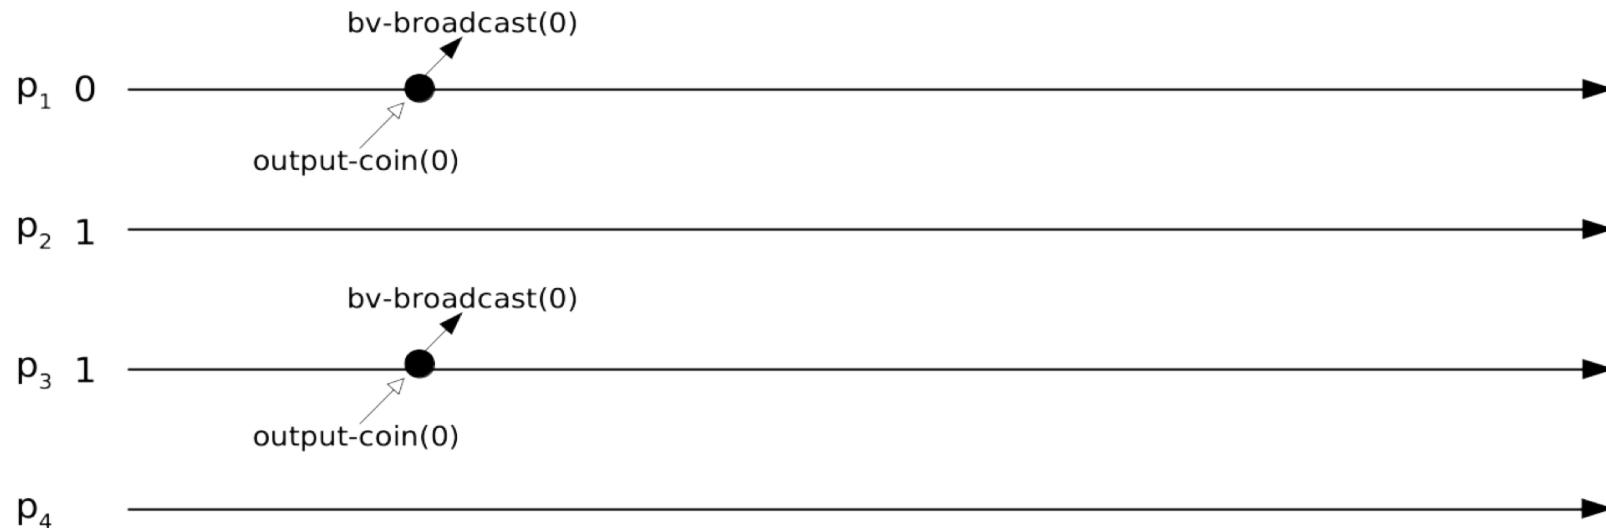
5: upon event release-coin do
6:   for all  $\mathcal{G} \in \mathcal{H}$  such that  $p_i \in \mathcal{G}$  do
7:     let  $s_{i\mathcal{G}}$  be the share of  $p_i$  for guild  $\mathcal{G}$ 
8:     for all  $p_j \in \mathcal{P}$  do
9:       send message [SHARE,  $s_{i\mathcal{G}}$ ,  $\mathcal{G}$ , round] to  $p_j$ 

10: upon receiving a message [SHARE,  $s$ ,  $\mathcal{G}$ , r] from  $p_j$  such that
11:    $r = \text{round}$  and  $p_j \in \mathcal{G}$  do
12:     if  $share[\mathcal{G}][j] = \perp$  then
13:        $share[\mathcal{G}][j] \leftarrow s$ 

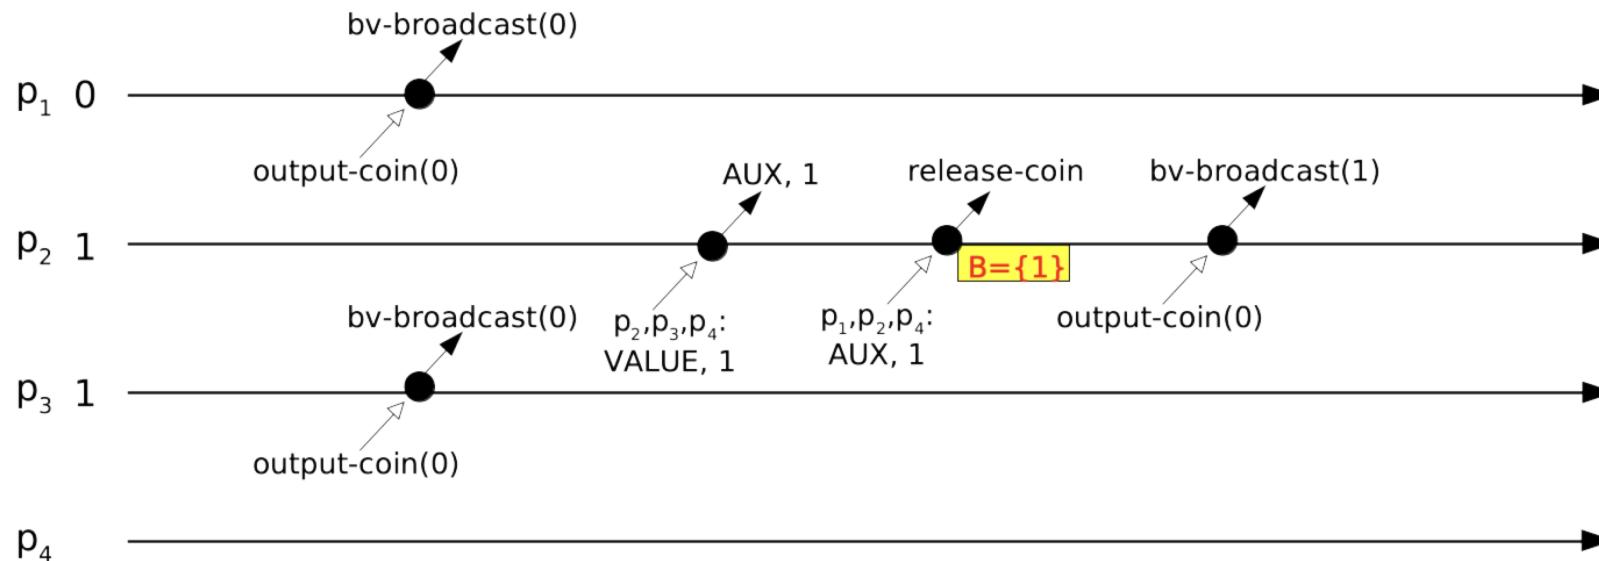
14: upon exists  $\mathcal{G}$  such that for all  $j$  with  $p_j \in \mathcal{G}$ , it holds
15:    $share[\mathcal{G}][j] \neq \perp$  do
16:      $s \leftarrow \sum_{j:p_j \in \mathcal{G}} share[\mathcal{G}][j]$ 
17:     output output-coin( $s$ )
```







Random coin = 0



Random coin = 0

